

# **CS5226 Lecture 6**

## **Transaction Tuning**

# Transaction

- ▶ **Transaction (Xact)** - any one execution of a user program
  - ▶ an abstraction representing a logical unit of work.
- ▶ A Xact  $T_i$  can be viewed as a sequence of **actions**:
  - ▶  $R_i(O) = T_i$  reads an object  $O$
  - ▶  $W_i(O) = T_i$  writes an object  $O$
  - ▶  $Commit_i = T_i$  completes successfully
  - ▶  $Abort_i = T_i$  terminates unsuccessfully and undoes all previous actions
- ▶ Each Xact must end with either a commit or an abort action

# Transaction (cont.)

- ▶ **Example:** A Xact  $T_1$  that transfers \$100 from  $A$  to  $B$  can be abstracted as

$T_1: R_1(A), W_1(A), R_1(B), W_1(B), Commit_1$

- ▶ **Xact consistency property:** If a Xact starts with a consistent DB state, then it also ends with a consistent DB state
- ▶ But, an intermediate DB state is not necessarily consistent!

# Transaction Schedule

- ▶ **Schedule** = a list of actions from a set of Xacts, where the order of the actions within each Xact is preserved
- ▶ **Example:** Consider two Xacts  $T_1$  and  $T_2$ :
  - ▶  $T_1$  transfers \$100 from  $A$  to  $B$   
 $T_1: R_1(A), W_1(A), R_1(B), W_1(B), Commit_1$
  - ▶  $T_2$  increments both  $A$  and  $B$  by 10%  
 $T_2: R_2(A), W_2(A), R_2(B), W_2(B), Commit_2$
- ▶ Some schedules of  $T_1$  and  $T_2$ :
  - $S_1: R_1(A), W_1(A), R_1(B), W_1(B), Commit_1, R_2(A), W_2(A), R_2(B), W_2(B), Commit_2$
  - $S_2: R_2(A), W_2(A), R_2(B), W_2(B), Commit_2, R_1(A), W_1(A), R_1(B), W_1(B), Commit_1$
  - $S_3: R_1(A), W_1(A), R_2(A), W_2(A), R_1(B), W_1(B), Commit_1, R_2(B), W_2(B), Commit_2$
- ▶ A **serial schedule** is a schedule where the actions of Xacts are not interleaved

# Read/Write From & Final Write

- ▶ We say that  $T_j$  reads  $O$  from  $T_i$  in a schedule  $S$  if the last write action on  $O$  before  $R_j(O)$  in  $S$  is  $W_i(O)$
- ▶ We say that  $T_i$  performs the final write on  $O$  in a schedule  $S$  if the last write action on  $O$  in  $S$  is  $W_i(O)$
- ▶ **Example:** Consider the following schedule.

$R_1(A), W_1(A), R_2(A), W_2(A), R_1(B), W_1(B), \text{Commit}_1, R_2(B), W_2(B), \text{Commit}_2$

- ▶  $T_2$  reads  $A$  from  $T_1$
- ▶  $T_2$  reads  $B$  from  $T_1$
- ▶  $T_2$  performs the final write on  $A$
- ▶  $T_2$  performs the final write on  $B$

# Interleaved Transaction Execution

- **Example:** Consider the Xacts  $T_1$  &  $T_2$  and four possible schedules:

$T_1$ : transfers \$100 from  $A$  to  $B$

$T_2$ : increments both  $A$  and  $B$  by 10%

$S_1$ :  $R_1(A), W_1(A), R_1(B), W_1(B), Commit_1, R_2(A), W_2(A), R_2(B), W_2(B), Commit_2$

$S_2$ :  $R_2(A), W_2(A), R_2(B), W_2(B), Commit_2, R_1(A), W_1(A), R_1(B), W_1(B), Commit_1$

$S_3$ :  $R_1(A), W_1(A), R_2(A), W_2(A), R_1(B), W_1(B), Commit_1, R_2(B), W_2(B), Commit_2$

$S_4$ :  $R_1(A), W_1(A), R_2(A), W_2(A), R_2(B), W_2(B), Commit_2, R_1(B), W_1(B), Commit_1$

If  $A = \$300$  and  $B = \$500$  then

- $S_1$  results in  $A = \$220$  and  $B = \$660$
- $S_2$  results in  $A = \$230$  and  $B = \$650$
- $S_3$  results in  $A = \$220$  and  $B = \$660$
- $S_4$  results in  $A = \$220$  and  $B = \$650$

# Anomalies with Interleaved Executions

- ▶ Two actions on the same object **conflict** if (1) at least one of them is a write action and (2) the actions are from different Xacts
- ▶ Anomalies can arise due to conflicting actions:
  1. **Dirty read problem (due to WR conflicts)**
    - ★  $T_2$  reads an object that has been modified by  $T_1$  (which has not yet committed)
    - ★  $T_2$  could see an inconsistent DB state!
  2. **Unrepeatable read problem (due to RW conflicts)**
    - ★  $T_2$  updates an object that  $T_1$  has just read while  $T_1$  is still in progress
    - ★  $T_1$  could get a different value if it reads the object again!
  3. **Lost update problem (due to WW conflicts)**
    - ★  $T_2$  overwrites the value of an object that has been modified by  $T_1$  while  $T_1$  is still in progress
    - ★  $T_1$ 's update is lost!

# Dirty Read Problem: Example

| $T_1$        | $T_2$            | Comments  |
|--------------|------------------|-----------|
|              |                  | $x = 100$ |
| R(x)         |                  | 100       |
| $x = x + 20$ |                  |           |
| W(x)         |                  | $x = 120$ |
|              | R(x)             | 120       |
|              | $x = x \times 2$ |           |
| Rollback     |                  |           |
|              | W(x)             | $x = 240$ |

- For every serial schedule, the final value of  $x$  is 200



# Unrepeatable Read Problem: Example

| $T_1$  | $T_2$        | Comments  |
|--------|--------------|-----------|
|        |              | $x = 100$ |
| $R(x)$ |              | 100       |
|        | $R(x)$       | 100       |
|        | $x = x - 20$ |           |
|        | $W(x)$       | $x = 80$  |
| $R(x)$ |              | 80        |

- For every serial schedule, both values read by  $T_1$  are the same

# Lost Update Problem: Example

| $T_1$        | $T_2$            | Comments  |
|--------------|------------------|-----------|
|              |                  | $x = 100$ |
| R(x)         |                  | 100       |
|              | R(x)             | 100       |
| $x = x + 20$ |                  |           |
|              | $x = x \times 2$ |           |
| W(x)         |                  | $x = 120$ |
|              | W(x)             | $x = 200$ |

- ▶ For schedule  $(T_1, T_2)$ , the final value of  $x$  is 240
- ▶ For schedule  $(T_2, T_1)$ , the final value of  $x$  is 220

# Conflict Serializable Schedule

- ▶ Two schedules  $S$  and  $S'$  are said to be **conflict equivalent** if
  1. they involve the same set of actions of the same Xacts, and
  2. they order every pair of conflicting actions of two committed Xacts in the same way

- ▶ **Example:** Consider the two schedules  $S$  and  $S'$ :

$S:$     .....  $W_i(X)$  .....  $R_j(X)$  .....

$S':$     .....  $R_j(X)$  .....  $W_i(X)$  .....

$S$  is not conflict equivalent to  $S'$

- ▶ A schedule is a **conflict serializable schedule** if it is conflict equivalent to a serial schedule

# Conflict Serializable Schedule: Example

- ▶ **Example:** Consider the following transactions and schedule:

$T_1$ :  $R_1(A)$ ,  $W_1(A)$ ,  $Commit_1$ ,

$T_2$ :  $W_2(A)$ ,  $Commit_2$

$T_3$ :  $W_3(A)$ ,  $Commit_3$


S:  $R_1(A)$ ,  $W_2(A)$ ,  $Commit_2$ ,  $W_1(A)$ ,  $Commit_1$ ,  $W_3(A)$ ,  $Commit_3$

- ▶ S is not a conflict serializable schedule
  - ▶ In any serial schedule  $S'$  with  $T_1$  preceding  $T_2$ , the orderings of  $W_1(A)$  and  $W_2(A)$  are different in S &  $S'$
  - ▶ In any serial schedule  $S'$  with  $T_2$  preceding  $T_1$ , the orderings of  $R_1(A)$  and  $W_2(A)$  are different in S &  $S'$

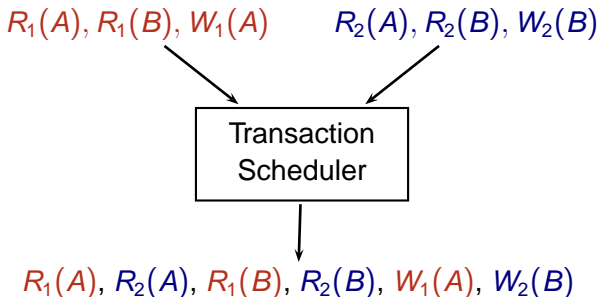
# Testing for Conflict Serializability

- ▶ A **conflict serializability graph** for a schedule  $S$  is a directed graph  $G = (V, E)$  such that
  - ▶  $V$  contains a node for each committed Xact in  $S$
  - ▶  $E$  contains  $(T_i, T_j)$  if an action in  $T_i$  precedes and conflicts with one of  $T_j$ 's actions
- ▶ **Theorem:** A schedule is conflict serializable iff its conflict serializability graph is acyclic
- ▶ **Example:** Conflict serializability graph for schedule

$R_1(A), W_2(A), \text{Commit}_2, W_1(A), \text{Commit}_1, W_3(A), \text{Commit}_3$



# Transaction Scheduler



- ▶ For each input action (read, write, commit, rollback) to the scheduler, the scheduler performs one of the following:
  - ▶ output the action to the schedule,
  - ▶ reject the action and abort the transaction, or
  - ▶ postpone the action by blocking the transaction

# Lock-Based Concurrency Control

- ▶ Each Xact needs to request for an appropriate **lock** on an object before the Xact can access the object
- ▶ **Locking modes**
  - ▶ **Shared (S) locks** for reading objects
  - ▶ **Exclusive (X) locks** for writing objects
- ▶ **Lock compatibility:**

| Lock Requested | Lock Held |   |   |
|----------------|-----------|---|---|
|                | -         | S | X |
| S              | ✓         | ✓ | × |
| X              | ✓         | × | × |

✓: Compatible  
Lock request is granted

×: Incompatible  
Lock request is blocked

- ▶ A Xact is **blocked** if its request for a lock on an object is incompatible with an existing lock on the object held by some other Xact.

# Lock-Based Concurrency Control (cont.)

## ► Notations

- $S_i(O)$ : Xact  $T_i$  is requesting S-lock on object O
- $X_i(O)$ : Xact  $T_i$  is requesting X-lock on object O
- $U_i(O)$ : Xact  $T_i$  releases lock on object O

## ► Example: $R_1(A)$ , $W_2(A)$ , $W_2(B)$ , $W_1(B)$

$S_1(A), R_1(A), U_1(A), X_2(A), W_2(A), X_2(B), W_2(B), U_2(A), U_2(B), X_1(B), W_1(B), U_1(B)$



# Two Phase Locking (2PL) Protocol

- ▶ **2PL Protocol:**

1. To read an object O, a Xact must hold a S-lock or X-lock on O
2. To write to an object O, a Xact must hold a X-lock on O
3. Once a Xact releases a lock, the Xact can't request any more locks

- ▶ Xacts using 2PL can be characterized into two phases:

- ▶ **Growing phase:** before releasing 1<sup>st</sup> lock
- ▶ **Shrinking phase:** after releasing 1<sup>st</sup> lock

**Theorem:** 2PL schedules are conflict serializable

# Two Phase Locking (2PL) Protocol (cont.)

- ▶ **Example:**  $R_1(A)$ ,  $W_2(A)$ ,  $W_2(B)$ ,  $W_1(B)$

$S_1(A)$ ,  $R_1(A)$ ,  $U_1(A)$ ,  $X_2(A)$ ,  $W_2(A)$ ,  $X_2(B)$ ,  $W_2(B)$ ,  $U_2(A)$ ,  $U_2(B)$ ,  $X_1(B)$ ,  $W_1(B)$ ,  $U_1(B)$

Not permitted by 2PL!



- ▶ The above example schedule is not a 2PL schedule

# Strict Two Phase Locking (strict 2PL) Protocol

- ▶ **Strict 2PL Protocol:**

1. To read an object O, a Xact must hold a S-lock or X-lock on O
2. To write to an object O, a Xact must hold a X-lock on O
3. A Xact must hold on to locks until Xact commits or aborts

- ▶ **Theorem:** strict 2PL schedules  $\subset$  2PL schedules

# Deadlocks

- ▶ **Deadlock:** cycle of Xacts waiting for locks to be released by each other
- ▶ **Example:**
  - $T_1$  requests X-lock on A and is granted;
  - $T_2$  requests X-lock on B and is granted;
  - $T_1$  requests X-lock on B and is blocked;
  - $T_2$  requests X-lock on A and is blocked;

# Phantom Read Phenomenon

Accounts

| account | branch | balance |
|---------|--------|---------|
| 101     | north  | 200     |
| 250     | south  | 1000    |
| 444     | south  | 800     |

Assets

| branch | total |
|--------|-------|
| north  | 200   |
| south  | 1800  |

| Transaction 1  | Transaction 2   |
|--|---|
| <b>select</b> sum (balance)<br><b>from</b> Accounts<br><b>where</b> branch = 'south';<br><br><b>select</b> total<br><b>from</b> Assets<br><b>where</b> branch = 'south'; | <b>insert into</b> Accounts<br><b>values</b> (222,'south',100);<br><br><b>update</b> Assets<br><b>set</b> total = total + 100<br><b>where</b> branch = 'south'; |

# Phantom Read Phenomenon (cont.)

- ▶ A transaction re-executes a query returning a set of rows that satisfy a search condition and finds that the set of rows satisfying the condition has changed due to another recently committed transaction.

```
select    sum (balance)
from      Accounts
where     branch = 'south';
```

```
insert into Accounts
           values (222,'south',100);
```

```
update Assets
      set total = total + 100
      where branch = 'south';
```

```
select    total
from      Assets
where     branch = 'south';
```

- ▶ Can be prevented by **predicate locking** via **index locking**

# ANSI SQL Isolation Levels

| Isolation Level  | Dirty Read   | Unrepeatable Read | Phantom Read |
|------------------|--------------|-------------------|--------------|
| READ UNCOMMITTED | possible     | possible          | possible     |
| READ COMMITTED   | not possible | possible          | possible     |
| REPEATABLE READ  | not possible | not possible      | possible     |
| SERIALIZABLE     | not possible | not possible      | not possible |

| Degree | Isolation level  | Write Locks   | Read Locks     | Phantom Read |
|--------|------------------|---------------|----------------|--------------|
| 0      | Read Uncommitted | long duration | none           | possible     |
| 1      | Read Committed   | long duration | short duration | possible     |
| 2      | Repeatable Read  | long duration | long duration  | possible     |
| 3      | Serializable     | long duration | long duration  | not possible |

# Transaction Chopping

## Transaction $T_{1,1}$

```
BEGIN TRANSACTION;  
SQL statements1;  
END TRANSACTION;
```

## Transaction $T_1$

```
BEGIN TRANSACTION;  
SQL statements1;  
SQL statements2;  
SQL statements3;  
END TRANSACTION;
```

## Transaction $T_{1,2}$

```
BEGIN TRANSACTION;  
SQL statements2;  
END TRANSACTION;
```

## Transaction $T_{1,3}$

```
BEGIN TRANSACTION;  
SQL statements3;  
END TRANSACTION;
```



# Transaction Chopping (cont.)

- ▶ A **chopping** partitions  $T$  into  $k$  pieces,  $k \geq 1$
- ▶ Each database access performed by  $T$  is in exactly one piece
- ▶ **Example 1:** Consider 2 transactions:
  - ▶  $T_1: R_1(a), W_1(a), R_1(y), W_1(y)$
  - ▶  $T_2: R_2(a)$
- ▶  $T_1$  could be chopped into 2 pieces:
  - ▶  $T_{1,1}: R_1(a), W_1(a)$                        $T_{1,2}: R_1(y), W_1(y)$

# Why Transaction Chopping?

- ▶ Transaction chopping can increase concurrency
- ▶ Possible S2PL schedules for  $\{T_1, T_2\}$ :
  - S1.  $R_2(a)$ ,  $R_1(a)$ ,  $W_1(a)$ ,  $R_1(y)$ ,  $W_1(y)$
  - S2.  $R_1(a)$ ,  $R_2(a)$ ,  $W_1(a)$ ,  $R_1(y)$ ,  $W_1(y)$
  - S3.  $R_1(a)$ ,  $W_1(a)$ ,  $R_1(y)$ ,  $W_1(y)$ ,  $R_2(a)$
- ▶ Possible S2PL schedules for  $\{T_{1,1}, T_{1,2}, T_2\}$ 
  - S1.  $R_2(a)$ ,  $R_1(a)$ ,  $W_1(a)$ ,  $R_1(y)$ ,  $W_1(y)$
  - S2.  $R_1(a)$ ,  $R_2(a)$ ,  $W_1(a)$ ,  $R_1(y)$ ,  $W_1(y)$
  - S3.  $R_1(a)$ ,  $W_1(a)$ ,  $R_1(y)$ ,  $W_1(y)$ ,  $R_2(a)$
  - S4.  $R_1(a)$ ,  $W_1(a)$ ,  $R_2(a)$ ,  $R_1(y)$ ,  $W_1(y)$
  - S5.  $R_1(a)$ ,  $W_1(a)$ ,  $R_1(y)$ ,  $R_2(a)$ ,  $W_1(y)$

Each of the 5 schedules is a conflict serializable schedule for  $\{T_1, T_2\}$

# CSG(S,T)

- ▶ Let  $T$  denote a set of transactions
- ▶ Let  $T'$  denote a chopping of  $T$
- ▶ Let  $S$  denote a schedule for  $T'$
- ▶ We use  $\text{CSG}(S, T')$  to denote the conflict serializability graph for schedule  $S$  wrt the set of transactions  $T'$
- ▶ If  $S$  is a S2PL schedule for  $T'$ , then
  - ▶  $\text{CSG}(S, T')$  must be acyclic
  - ▶ But  $\text{CSG}(S, T)$  could be cyclic

# Incorrect Transaction Chopping

- ▶ **Example 2:** Consider 3 transactions:

- ▶  $T_1$ :  $R_1(a)$ ,  $W_1(a)$ ,  $R_1(y)$ ,  $W_1(y)$
- ▶  $T_2$ :  $R_2(b)$ ,  $W_2(b)$ ,  $R_2(y)$ ,  $W_2(y)$
- ▶  $T_3$ :  $R_3(a)$ ,  $R_3(b)$ ,  $R_3(y)$

- ▶ Suppose each of  $T_1$  &  $T_2$  is chopped into 2 pieces:

- ▶  $T_{1,1}$ :  $R_1(a)$ ,  $W_1(a)$                        $T_{1,2}$ :  $R_1(y)$ ,  $W_1(y)$
- ▶  $T_{2,1}$ :  $R_2(b)$ ,  $W_2(b)$                        $T_{2,2}$ :  $R_2(y)$ ,  $W_2(y)$

- ▶ A possible S2PL schedule for  $\{T_{1,1}, T_{1,2}, T_{2,1}, T_{2,2}, T_3\}$ :


$R_1(a)$ ,  $W_1(a)$ ,  $R_1(y)$ ,  $W_1(y)$ ,  $R_2(b)$ ,  $W_2(b)$ ,  $R_3(a)$ ,  $R_3(b)$ ,  $R_3(y)$ ,  $R_2(y)$ ,  $W_2(y)$

Not a conflict serializable schedule for  $\{T_1, T_2, T_3\}$ !


# Incorrect Transaction Chopping (cont.)

- ▶  $T = \{T_1, T_2, T_3\}$
- ▶  $T' = \{T_{1,1}, T_{1,2}, T_{2,1}, T_{2,2}, T_3\}$
- ▶ A S2PL schedule  $S$  for  $T'$ :

$R_1(a), W_1(a), R_1(y), W_1(y), R_2(b), W_2(b), R_3(a), R_3(b), R_3(y), R_2(y), W_2(y)$



**CSG(S, T)**



# Rollback-Safe Chopping

## Transaction $T_1$

```
BEGIN TRANSACTION;  
SQL statements1;  
if (condition1) then  
    ROLLBACK;  
SQL statements2;  
if (condition2) then  
    ROLLBACK;  
SQL statements3;  
END TRANSACTION;
```

## Transaction $T_{1,1}$

```
BEGIN TRANSACTION;  
SQL statements1;  
if (condition1) then ROLLBACK;  
END TRANSACTION;
```

## Transaction $T_{1,2}$

```
BEGIN TRANSACTION;  
SQL statements2;  
if (condition2) then ROLLBACK;  
END TRANSACTION;
```

## Transaction $T_{1,3}$

```
BEGIN TRANSACTION;  
SQL statements3;  
END TRANSACTION;
```

# Rollback-Safe Chopping (cont.)

## Transaction $T_1$

```
BEGIN TRANSACTION;  
SQL statements1;  
if (condition1) then  
    ROLLBACK;  
SQL statements2;  
if (condition2) then  
    ROLLBACK;  
SQL statements3;  
END TRANSACTION;
```

## Transaction $T_{1,1}$

```
BEGIN TRANSACTION;  
SQL statements1;  
if (condition1) then ROLLBACK;  
SQL statements2;  
if (condition2) then ROLLBACK;  
END TRANSACTION;
```

## Transaction $T_{1,2}$

```
BEGIN TRANSACTION;  
SQL statements3;  
END TRANSACTION;
```

# Rollback-Safe Chopping (cont.)

- ▶ A chopping of T is **rollback-safe** if
  - ▶ either T has no rollback statements
  - ▶ or all rollback statements of T are in its first piece
- ▶ All the statements in the first piece must execute before any other statement of T
- ▶ A chopping of a set of transactions is **rollback-safe** if the chopping of each of its transaction is rollback-safe



# Execution rules

1. Each piece acts like a transaction: follows strict 2PL protocol
2. Order of execution of pieces follow order of transaction
3. If a piece is aborted because of a lock conflict, then it will be resubmitted repeatedly until it commits.
4. If a piece is aborted because of a rollback, then no other pieces for that transaction will execute.

# Correct Chopping

- ▶ A chopping of  $\{T_1, \dots, T_n\}$  is **correct** if **any** execution of the chopping that obeys the execution rules is *conflict equivalent* to some serial execution of the original transactions.

# Conflicting Pieces

Two pieces  $p$  and  $p'$  are **conflicting** if all the following conditions hold:


1.  $p$  &  $p'$  belong to different transactions,
2.  $p$  contains a step that writes an object  $O$ , and
3.  $p'$  contains a step that reads/writes the same object  $O$

# Chopping Graph

- ▶ Undirected graph
- ▶ Each vertex represents a chopped piece
- ▶ Two types of edges:
  - ▶ conflict edges (C-edges)
  - ▶ sibling edges (S-edges)
- ▶  $(p, p')$  is a **S-edge** if  $p$  &  $p'$  belong to the same transaction
- ▶  $(p, p')$  is a **C-edge** if  $p$  &  $p'$  is a pair of conflicting pieces

## Example 3

- ▶  $T_1: R_1(x), W_1(x), R_1(y), W_1(y)$
- ▶  $T_2: R_2(x), W_2(x)$
- ▶  $T_3: R_3(y), W_3(y)$
- ▶ Suppose that  $T_1$  is chopped into two pieces:
  - ▶  $T_{1,1}: R_1(x), W_1(x)$
  - ▶  $T_{1,2}: R_1(y), W_1(y)$




S-edge: =

C-edge: —

# Review: Graph Concepts

- ▶ A **simple cycle** consists of a sequence of distinct nodes  $\langle v_1, v_2, \dots, v_n \rangle$  such that
  1. there is an edge  $(v_i, v_{i+1})$  for each  $i \in [1, n)$ , and
  2. there is an edge  $(v_n, v_1)$
- ▶ A **connected component**  $C$  in a graph  $G$  is a maximal subgraph of  $G$  such that for every pair of nodes  $v$  &  $v'$  in  $C$ ,  $v$  is connected to  $v'$  by some path




# SC-cycle

- ▶ A chopping graph has a **SC-cycle** if it contains a simple cycle that includes at least one S-edge and at least one C-edge

# Revisiting Example 3

- ▶  $T_1: R_1(x), W_1(x), R_1(y), W_1(y)$
- ▶  $T_2: R_2(x), W_2(x)$
- ▶  $T_3: R_3(y), W_3(y)$
- ▶ Suppose that  $T_1$  is chopped into two pieces:
  - ▶  $T_{1,1}: R_1(x), W_1(x)$
  - ▶  $T_{1,2}: R_1(y), W_1(y)$




S-edge: =  
C-edge: —



## Revisiting Example 3 (cont.)

Suppose that  $T_{1,1}$  is further chopped into two pieces:

- ▶  $T_{1,1,1}$ :  $R_1(x)$
- ▶  $T_{1,1,2}$ :  $W_1(x)$



- ▶  $T_{1,2}$ :  $R_1(y)$ ,  $W_1(y)$
- ▶  $T_2$ :  $R_2(x)$ ,  $W_2(x)$
- ▶  $T_3$ :  $R_3(y)$ ,  $W_3(y)$

# Correct Chopping

**Theorem 1:** A chopping is correct if

1. it is rollback-safe, and
2. its chopping graph contains no SC-cycle

# Quiz 1

- ▶ Consider the following transactions
  - ▶  $T_1: R_1(x), W_1(x), R_1(y), W_1(y)$
  - ▶  $T_2: R_2(x)$
  - ▶  $T_3: R_3(y), W_3(y)$
- ▶ Suppose that  $T_1$  is chopped into 3 pieces:
  - ▶  $T_{1,1}: R_1(x)$
  - ▶  $T_{1,2}: W_1(x)$
  - ▶  $T_{1,3}: R_1(y), W_1(y)$
- ▶ Is this chopping correct?

# Properties of Chopping

**Lemma 1:** If a chopping graph contains an SC-cycle, then any further chopping of any of the transactions will not render it acyclic

**Lemma 2:** If two pieces of transaction  $T$  are in an SC-cycle as the result of some chopping, then they will be in a cycle even if no other transactions are chopped.

# Private chopping

- ▶ Consider a set of transactions  $T = \{T_1, \dots, T_n\}$  that run at an interval
- ▶ Let  $\text{chop}(T_i)$  denote a chopping of  $T_i$
- ▶  $\text{chop}(T_i)$  is a **private chopping** of  $T_i$ , denoted by  $\text{private}(T_i)$ , if
  1.  $\text{chop}(T_i)$  is a rollback-safe chopping of  $T_i$ , and
  2. there is no SC-cycle in the graph whose nodes are  $(T - \{T_i\}) \cup \text{chop}(T_i)$

**Theorem 2:**

The chopping consisting of  $\{\text{private}(T_1), \dots, \text{private}(T_n)\}$  is rollback safe and has no SC-cycles.

# Finest chopping

- ▶ Let **FineChop**( $T_i$ ) denote the **finest chopping** of  $T_i$  with the following two properties:
  1. FineChop( $T_i$ ) is a private chopping of  $T_i$
  2. If piece  $p$  is a member of FineChop( $T_i$ ), then there is no other private chopping of  $T_i$  containing  $p_1$  and  $p_2$  such that  $p_1$  and  $p_2$  partition  $p$  and neither is empty

The finest chopping of a set of transactions  $S = \{T_1, \dots, T_n\}$  is given by  $\{ \text{FineChop}(T_1), \dots, \text{FineChop}(T_n) \}$

# Finest chopping: algorithm

**Input:** A set of transactions  $S$  & a transaction  $T \in S$


**Output:** FineChop( $T$ )

01. if there are rollback statements in  $T$  then
02.      $p_1 \leftarrow$  all database writes of  $T$  that may occur before  
          or concurrently with any rollback statement of  $T$
03. else
04.      $p_1 \leftarrow$  set consisting of the first database access
05.  $P \leftarrow \{\{x\} \mid x \text{ is a database access not in } p_1\}$
06.  $P \leftarrow P \cup \{p_1\}$
07. let  $G = (V, E)$  be an undirected graph, where  
       $V = (S - \{T\}) \cup P$ ,  $E$  is the set of C-edges among  $V$
08. for each connected component  $C_i$  in  $G$  do
09.     merge all pieces of  $P$  in  $C_i$  into a single piece
10. return the set of pieces in  $P$


# Revisiting Example 3

- ▶  $T_1: R_1(x), W_1(x), R_1(y), W_1(y)$
- ▶  $T_2: R_2(x), W_2(x)$
- ▶  $T_3: R_3(y), W_3(y)$


$$\text{FineChop}(T_1) = \{\{R_1(x), W_1(x)\}, \{R_1(y), W_1(y)\}\}$$



$$\text{FineChop}(T_2) = \{\{R_2(x), W_2(x)\}\}$$



$$\text{FineChop}(T_3) = \{\{R_3(y), W_3(y)\}\}$$





# Quiz 2

Find the finest chopping for the set of 4 transactions:

- ▶  $T_1: R_1(a), W_1(a), R_1(x), W_1(x)$
- ▶  $T_2: R_2(c), W_2(c), R_2(x), W_2(x)$
- ▶  $T_3: R_3(d), W_3(d), R_3(y), W_3(y)$
- ▶  $T_4: R_4(a), R_4(b), R_4(c), R_4(x), R_4(d), R_4(e), R_4(y)$

# References

## Required Readings

- ▶ Section 2.2, Locking and Concurrency Control, Shasha & Bonnet's book
- ▶ Appendix B, Transaction Chopping, Shasha & Bonnet's book

## Additional Readings

- ▶ D. Shasha, et al., Transaction Chopping: Algorithms and Performance Studies, ACM TODS 20(3), 325-363, 1995