

CS5226 Lecture 9


Statistics Tuning

Statistics Tuning Issues

- ▶ Which statistics to collect?
- ▶ At what level of detail of the statistics to collect?
- ▶ When to collect/refresh?
- ▶ Challenges
 - ▶ outdated statistics
 - ▶ improperly configured statistics

Cost Estimation of Query Plan

- ▶ Cost estimation involves the following:
 1. What is the evaluation cost of each operation?
 - ★ Depends on: size of input operands, available buffer pages, available indexes, etc.
 2. What is the output size of each operation?



Examples of Statistics

- ▶ cardinality (i.e. number of tuples)
- ▶ number of distinct values in each column
- ▶ the highest & lowest values in each column
- ▶ column group statistics
- ▶ histograms
- ▶ frequent values of some columns

Size Estimation

- ▶ Consider a query $q = \sigma_p(e)$

- ▶ $p = t_1 \wedge t_2 \wedge \dots \wedge t_n$

- ▶ $e = R_1 \times R_2 \times \dots \times R_m$

- ▶ How to estimate $||q||$?

- ▶ We have $||e|| = \prod_{i=1}^m ||R_i|| = ||R_1|| \times ||R_2|| \times \dots \times ||R_m||$

- ▶ Each term t_i potentially eliminates some tuples in e
- ▶ **Reduction factor** of a term t_i (denoted by rf_i) is the fraction of tuples in e that satisfy t_i ; i.e., $rf_i = \frac{||\sigma_{t_i}(e)||}{||e||}$
- ▶ Assuming the terms in p are statistically independent,

$$||q|| = ||e|| \times \prod_{i=1}^n rf_i = ||e|| \times rf_1 \times rf_2 \times \dots \times rf_n$$

Estimation assumptions

- ▶ **Uniformity assumption**

- ▶ uniform distribution of attribute values

- ▶ **Independence assumption**

- ▶ Independent distribution of values in different attributes

- ▶ **Inclusion assumption**

- ▶ For a join predicate $R.A = S.B$, we have $(\pi_A(R) \subseteq \pi_B(S))$
or $(\pi_A(R) \supseteq \pi_B(S))$

How to estimate reduction factor?

Form of t_i	Estimation of rf_i
$A_j = v$	$\frac{1}{N_{key}(I)}$ if I_{A_j} exists, $\frac{1}{10}$ otherwise.
$A_j > v$	$\frac{I_{High}(I_{A_j}) - v}{I_{High}(I_{A_j}) - I_{Low}(I_{A_j})}$ if I_{A_j} exists, $\frac{1}{10}$ otherwise.
$A_j = A_k$	$\frac{1}{\max\{N_{key}(I_{A_j}), N_{key}(I_{A_k})\}}$ if both I_{A_j} and I_{A_k} exist, $\frac{1}{N_{key}(I_{A_j})}$ if only one index (say I_{A_j}) exists, $\frac{1}{10}$ otherwise.


- ▶ $N_{key}(I_{A_j})$ = number of keys in index I_{A_j}
- ▶ $I_{Low}(I_{A_j})$ = minimum key value in index I_{A_j}
- ▶ $I_{High}(I_{A_j})$ = maximum key value in index I_{A_j}

Improved Estimation using Histograms

- ▶ **histogram** = statistical information maintained by DBMS to estimate data distribution
- ▶ Main idea:
 - ▶ partition attribute's domain into sub-ranges called **buckets**
 - ▶ assume distribution within each bucket is uniform
- ▶ Types of histograms:
 - ▶ **Equiwidth histogram**: each bucket has equal number of values
 - ▶ **Equidepth histogram**: each bucket has equal number of tuples


Uniform vs Non-uniform Distributions

- ▶ Total number of distinct values = 15
- ▶ Total number of tuples = 45
- ▶ Uniform distribution assumption: each value has $\frac{45}{15} = 3$ tuples




Equiwidth vs Equidepth Histograms

Bucket No	Equiwidth		Equidepth	
	Value Range	No. Tuples	Value Range	No. Tuples
1	[0, 2]	$2+3+3=8$	[0, 3]	$2+3+3+1=9$
2	[3, 5]	$1+2+1=4$	[4, 6]	$2+1+8=11$
3	[6, 8]	$8+3+4=15$	[7, 9]	$3+4+2=9$
4	[9, 11]	$2+0+1=3$	[10, 13]	$0+1+2+4=7$
5	[12, 14]	$2+4+9=15$	[14, 14]	9



Equiwidth Histogram



Equidepth Histogram

Statistics Tuning Approaches

Reactive approaches

- ▶ Monitor a query during execution
- ▶ Observe errors between estimates and actual values from query feedback
- ▶ Possible strategies:
 - ▶ Use errors as adjustment factors to correct statistics for future queries,
 - ▶ Trigger statistics collection when error exceeds some threshold, or
 - ▶ Re-optimize current query

Statistics Tuning Approaches (cont.)

Proactive approaches

- ▶ Try to predict, identify and possibly solve potential problems before query execution
- ▶ Possible strategies:
 - ▶ Monitors update/delete/insert (UDI) activity on data & automatically refreshes statistics when UDI activity exceeds some threshold
 - ▶ Analyze queries to determine which statistics to collect
 - ▶ Maintains multiple plans for each query & adaptively change plan based on query runtime feedback

Statistical Views (Statviews)

- ▶ A **statview** is a view definition augmented with statistics collected on the result of executing this view, without the actual data
- ▶ DB2 Example:

```
create view sv_cust_order as (  
  select *  
  from Customer C, Order O  
  where C.cust# = O.cust#)
```

```
alter view sv_cust_order enable query optimization
```

```
runstats on table DB2DBA.sv_cust_order with distribution
```

Statistical Views (Statviews) (cont.)

Query: **select** *
from Customer C, Order O
where C.cust# = O.cust#
and C.country = 'Singapore'

- ▶ Without statview sv_cust_order,
 - ▶ estimate selectivity of selection predicate on C
 - ▶ estimate selectivity of join predicate on $C \times O$
 - ▶ apply independence assumption to estimate cardinality of result
- ▶ With statview sv_cust_order,
 - ▶ estimate selectivity of selection predicate on statview

Statview Tuning

- ▶ **Input:**

- ▶ a workload $W = \{Q_1, \dots, Q_n\}$, each Q_i is a query,
- ▶ c_{max} , the maximum number of statviews that can be maintained

- ▶ **Output:**

- ▶ A set of at most c_{max} statviews that minimizes the execution time of W

Statview Tuning (cont.)

- ▶ Let P denote a query plan
- ▶ C = **estimated cost of P** based on some simplifying assumptions (e.g. uniformity, independence, or inclusion)
- ▶ C_{acc}^P = **accurate cost estimate of P** where the cardinality at each operator in P is estimated without any simplifying assumptions
- ▶ Cost of P is **overestimated** if $C > C_{acc}^P$
- ▶ Cost of P is **underestimated** if $C < C_{acc}^P$

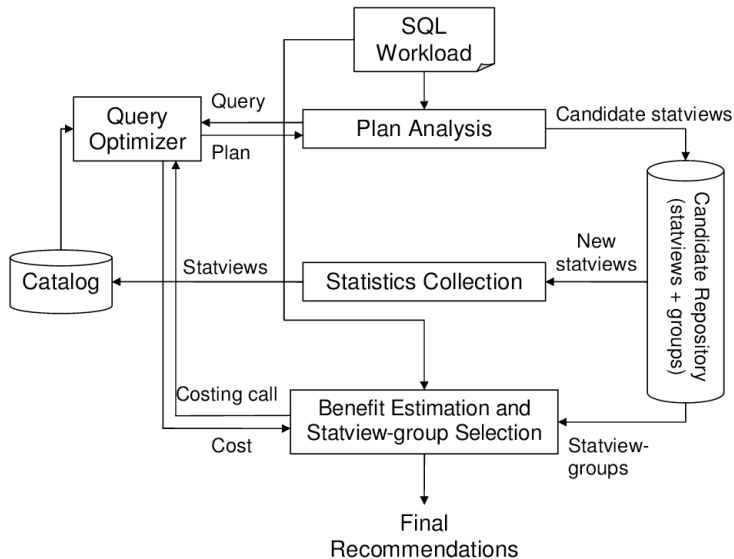
Relevant statviews

- ▶ E = subexpression in a plan P for query Q
- ▶ V = statview corresponding to E
- ▶ V is a **relevant statview** for optimizing Q if
 1. cost of E is **underestimated** & P is chosen by optimizer, or
 2. cost of E is **overestimated** & P will be chosen by optimizer when cost of E is corrected
- ▶ First type identified by “optimize & find-statviews” loop
- ▶ Second type identified by query structure analysis
 - ▶ predicates involving attributes with skewed distributions, arithmetic expressions, or UDFs
 - ▶ join predicates that violate inclusion assumption, etc

Statview Tuning

- ▶ **Optimization goal:**
 - ▶ Collect enough statviews to find optimal plan
- ▶ **Revised optimization goal:**
 - ▶ Collect enough statviews to find **predictable plan**
 - ▶ Plan cost is not underestimated

StatAdvisor Architecture



StatAdvisor

Input: Query workload W & c_{max} constraint


Output: Set of recommended statviews for W

01. $(V, G) \leftarrow \text{PlanAnalysis}(W)$
02. while $(V \neq \emptyset)$ do
03. $\text{CollectStatviews}(V)$
04. $(V, G) \leftarrow \text{PlanAnalysis}(W)$
05. $\text{EstimateBenefit}(G)$
06. $R \leftarrow \text{StatviewGroupSelection}(G, c_{max})$
07. return R

How to collect statistics for a statview?

- ▶ Statview V is on a **single table** R
 - ▶ Create a random sample S of R
 - ▶ Scan S to compute statistics for V
- ▶ Statview V is on **multiple joined tables** $R_1 \bowtie \dots \bowtie R_n$
 - ▶ Let R_1 be the root table
 - ★ R_1 has a foreign key to each of the tables R_2, \dots, R_n
 - ★ Each of the tables R_2, \dots, R_n has no foreign key to R_1
 - ▶ Compute the **join synopsis** J for the root table R_1
 1. Create a uniform random sample S of R_1
 2. For each table R_i that R_1 has a foreign key to, join S with R_i
 3. Repeat step 2 recursively
 - ▶ Scan J to compute statistics for V

Enumeration of candidate statviews



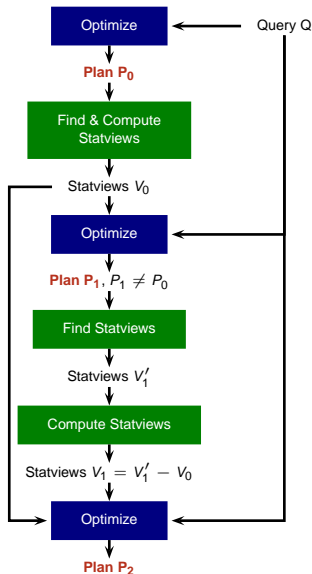
Case 1: $P_1 = P_0$

- ▶ V_0 is not beneficial for optimizing Q

Case 2: $P_1 \neq P_0$

- ▶ Either cost of P_1 is underestimated & P_0 could be a better plan than P_1
- ▶ Or P_1 is indeed a better plan than P_0

Enumeration of candidate statviews



Case 1: $P_2 = P_0$

- ▶ V_1 corrected the underestimation of cost of P_1

Case 2: $P_2 = P_1$

- ▶ V_1 is not beneficial for optimizing Q

Case 3: $(P_2 \neq P_0)$ and $(P_2 \neq P_1)$

- ▶ Either cost of P_2 is underestimated & P_0/P_1 could be a better plan than P_2
- ▶ Or P_2 is indeed a better plan than P_0 & P_1

Enumeration algorithm for query

Input: Query Q

Output: Candidate statviews for Q

01. $V \leftarrow \emptyset$; stop $\leftarrow false$; $k \leftarrow 0$;
02. repeat
03. generate the optimal plan P_k for Q using statviews V
04. find the important statviews VT for plan P_k
05. $VT \leftarrow VT - V$
06. if $(VT = \emptyset)$ or $(P_k = P_j \text{ for some } j < k)$ then
07. stop $\leftarrow true$
08. else
09. compute statviews VT
10. $V \leftarrow V \cup VT$
11. $k \leftarrow k + 1$
12. until (stop)
13. return V // V is a statview-group for Q

Enumeration algorithm for query (cont.)

- ▶ Let P be optimal plan for Q using V
 - ▶ V = the set of returned statviews
- ▶ P is guaranteed to be a predictable plan
- ▶ The approach could miss a better plan P' if
 1. the cost of some subexpression E of P' is overestimated, and
 2. E is not a subexpression of any enumerated plan

Enumeration algorithm for workload

Input: Workload $W = \{Q_1, \dots, Q_n\}$

Output: Candidate statviews for W

01. for each query $Q_i \in W$ do $V_i \leftarrow \emptyset$
02. $k \leftarrow 0$
03. repeat
04. for each query $Q_i \in W$ do
05. generate the optimal plan $P_{i,k}$ for each Q_i using V_i
06. find the important statviews VT_i for $P_{i,k}$
07. $VT_i \leftarrow VT_i - V_i$
08. if ($VT_i = \emptyset$) or ($P_{i,k} = P_{i,j}$ for some $j < k$) then
09. remove Q_i from W
10. else
11. $V_i \leftarrow V_i \cup VT_i$
12. partition $\{VT_j \mid Q_j \in W\}$ into batches $\{B_1, \dots, B_m\}$
13. compute statviews for each batch B_i
14. $k \leftarrow k + 1$
15. until (W is empty)
16. return ($V_1 \cup \dots \cup V_n$)

Benefit estimation

- ▶ P = optimal plan without statviews
- ▶ P_V = optimal plan with statviews V
- ▶ $\mathbf{B(V, Q)}$ = benefit of statviews V for query Q
 - ▶ $B(V, Q) = \text{ActCost}(Q, P) - \text{ActCost}(Q, P_V)$
 - ▶ $\text{ActCost}(Q, P)$ = actual execution cost of query Q using plan P
- ▶ $\mathbf{B'(V, Q)}$ = estimated benefit of statviews V for query Q
 - ▶ $B'(V, Q) = \text{EstCost}(P, V) - \text{EstCost}(P_V, V)$
 - ▶ $\text{EstCost}(P, V)$ = estimated execution cost of plan P with statviews V
 - ▶ Why not $\text{EstCost}(P, \emptyset) - \text{EstCost}(P_V, V)$?

Benefit estimation (cont.)

$\mathbf{B}(\mathbf{V}, \mathbf{W})$ = benefit of V for a workload W

- ▶ $W_V = \{Q \in W \mid Q \text{ generated } V\}$

- ▶ $B(V, W) = \sum_{Q \in W_V} B'(V, Q)$

Statview-group selection

Input: Set of statview-groups G & constraint c_{max}

Output: Set of recommended statview-groups R

```
01.  $c \leftarrow 0$ ;  $R \leftarrow \emptyset$ 
02. while ( $|G| > 0$ ) and ( $c < c_{max}$ ) do
03.      $V_{best} \leftarrow null$ 
04.      $B_{best} \leftarrow 0$ 
05.     for each  $V \in G$  do
06.         if ( $B(V) > B_{best}$ ) and ( $C(V, R) \leq c_{max} - c$ ) then
07.              $V_{best} \leftarrow V$ 
08.              $B_{best} \leftarrow B(V)$ 
09.     if ( $V_{best} \neq null$ ) then
10.          $c \leftarrow c + C(V_{best}, R)$ 
11.          $G \leftarrow G - V_{best}$ 
12.          $R \leftarrow R \cup V_{best}$ 
13.     else
14.         break
15. return  $R$ 
```

References

Required Readings

- ▶ A. El-Helw, I.F. Ihab, and C. Zuzarte, *StatAdvisor: recommending statistical views*, VLDB 2009

Additional Readings

- ▶ C. A. Galindo-Legaria, M. Joshi, F. Waas, and M.-C. Wu, *Statistics on Views*, VLDB 2003.
- ▶ A. Aboulnaga, P. Haas, M. Kandil, S. Lightstone, G. Lohman, V. Markl, I. Popivanov, V. Raman, *Automated statistics collection in DB2 UDB*, VLDB 2004.