

# **CS5226 Lecture 2**

## **Query Tuning**

# Query Tuning

- ▶ Eliminating redundant DISTINCT
- ▶ Rewriting nested queries

# Review: Relational Algebra

- ▶  $\pi_S R$  = projection of relation  $R$  onto the set of columns in  $S$
- ▶  $\pi_S^{all} R$  = projection of relation  $R$  onto the set of columns in  $S$  (duplicates preserved)
- ▶  $\sigma_p R$  = selection of relation  $R$ 
  - ▶  $p$  = predicate to filter qualifying rows from  $R$
- ▶  $\oplus$  = union with duplicates preserved (SQL's union all)
- ▶  $columns(X)$  = set of columns in  $X$ , where  $X$  is a relation or predicate

# Review: Relational Algebra (cont.)

- ▶  $R \times S$  = cross product
- ▶  $R \bowtie_p S$  = inner-join
- ▶  $R \ltimes_p S$  = left semi-join
- ▶  $R \rtimes_p S$  = right semi-join
- ▶  $R \rhd_p S$  = anti-join
- ▶  $R \Join_p S$  = left outer-join (LOJ)
- ▶  $R \Join_p S$  = right outer-join (ROJ)
- ▶  $R \Join_p S$  = full outer-join (FOJ)

# Review: Functional Dependencies

- ▶ **Functional dependencies (FDs)** are constraints on schemas that specify that the values for a certain set of attributes determine unique values for another set of attributes
- ▶ Let  $\alpha$  and  $\beta$  denote subsets of attributes of a relational schema  $R$  (i.e.,  $\alpha, \beta \subseteq \text{columns}(R)$ )
- ▶ We use  $\alpha \rightarrow \beta$  to denote that  $\alpha$  **functionally determines**  $\beta$  (or  $\beta$  functionally depends on  $\alpha$ )

# Review: Functional Dependencies (cont.)

- ▶ Let  $\pi_\alpha(t)$  denote the projection of  $\alpha$  on tuple  $t$
- ▶ Let  $r$  be a relation instance of relation schema  $R$
- ▶  $r$  satisfies FD  $\alpha \rightarrow \beta$  if for every pair of tuples  $t_1$  and  $t_2$  in  $r$  such that  $\pi_\alpha(t_1) = \pi_\alpha(t_2)$ , it is also true that  $\pi_\beta(t_1) = \pi_\beta(t_2)$
- ▶ A FD  $\alpha \rightarrow \beta$  holds on  $R$  iff for any relation instance  $r$  of  $R$ ,  $r$  satisfies  $\alpha \rightarrow \beta$

# Review: Functional Dependencies (cont.)

- ▶ Consider the following relation instance  $r$ :

Module	Prof	Room	Building	Time
CS101	Turing	LT 1	CS	0800
CS400	Turing	LT 1	CS	1400
MU300	Bach	LT 2	Math	1400
MA200	Newton	LT 2	Math	1000
CS101	Turing	LT 2	Math	1200

- ▶  $r$  satisfies  $\text{Room} \rightarrow \text{Building}$
- ▶  $r$  does not satisfy  $\text{Prof} \rightarrow \text{Module}$

# Review: Functional Dependencies (cont.)

- ▶ A set of attributes  $\alpha$  is a **superkey** of schema  $R$  if for every instance  $r$  of  $R$ ,  $r$  has no duplicate values for  $\alpha$
- ▶  $\alpha$  is a superkey of  $R$  iff  $\alpha \rightarrow \text{columns}(R)$
- ▶  $\alpha$  is a **key** of  $R$  if  $\alpha$  is a superkey of  $R$  and no proper subset of  $\alpha$  is a superkey of  $R$



# Review: Reasoning about FDs

- ▶ Let  $F$  be a set of FDs and  $f$  be an FD
- ▶  $F$  **logically implies (or implies)**  $f$  if every relation instance of  $R$  that satisfies the FDs  $F$  also satisfies the FD  $f$

## Inference Rules for FDs

- ▶ Let  $\alpha, \beta, \gamma \subseteq \text{columns}(R)$
- ▶ **Reflexivity**: If  $\beta \subseteq \alpha$ , then  $\alpha \rightarrow \beta$
- ▶ **Augmentation**: If  $\alpha \rightarrow \beta$ , then  $\alpha\gamma \rightarrow \beta\gamma$
- ▶ **Transitivity**: If  $\alpha \rightarrow \beta$  and  $\beta \rightarrow \gamma$ , then  $\alpha \rightarrow \gamma$
- ▶ **Union**: If  $\alpha \rightarrow \beta$  and  $\alpha \rightarrow \gamma$ , then  $\alpha \rightarrow \beta\gamma$
- ▶ **Decomposition**: If  $\alpha \rightarrow \beta$ , then  $\alpha \rightarrow \beta'$  for any  $\beta' \subseteq \beta$

# Review: Reasoning about FDs (cont.)

- ▶ Let  $\alpha \subseteq \text{columns}(R)$
- ▶ Let  $F$  be a set of FDs that hold on  $R$
- ▶ The **closure of  $\alpha$**  (with respect to  $F$ ), denoted by  $\alpha^+$ , is the set of attributes that are functionally determined by  $\alpha$  with respect to  $F$

$$\alpha^+ = \{A \in \text{columns}(R) \mid F \text{ implies } \alpha \rightarrow A\}$$

- ▶  $F$  implies  $\alpha \rightarrow \beta$  if and only if  $\beta \subseteq \alpha^+$  (w.r.t.  $F$ )

# Review: Reasoning about FDs (cont.)

**Input:**  $\alpha, F$

**Output:**  $\alpha^+$  (w.r.t.  $F$ )

$\alpha^+ = \alpha$

stop = false

repeat

if (there exists some FD  $\beta \rightarrow \gamma \in F$  such that  $\beta \subseteq \alpha^+$  and  $\gamma \notin \alpha^+$ ) then

$\alpha^+ = \alpha^+ \cup \gamma$

else

stop = true

until (stop)

return  $\alpha^+$

# Eliminating Redundant DISTINCT

- ▶ Customer (cust#, cname, country)
- ▶ SalesRep (rep#, sname, country)
- ▶ SalesOffice (country, address, phone)

**Q1:**     **select distinct** cust#  
         **from**     Customer, SalesRep  
         **where**    Customer.country = SalesRep.country

**Q2:**     **select distinct** cust#  
         **from**     Customer, SalesOffice  
         **where**    Customer.country = SalesOffice.country

# Eliminating Redundant DISTINCT

- ▶ **R(A,B,C)**
- ▶ **S(D,E,F)**

1. **select distinct B from R**
2. **select distinct A, B from R**
3. **select distinct A, D, E from R, S**
4. **select distinct A, E from R, S**
5. **select distinct A, E from R, S where A = F**
6. **select distinct A, E from R, S where B = D**

# SPJ Queries

A **SPJ query** is a query of the form

$$\pi_S^{all} \sigma_p(R_1 \times \cdots \times R_n)$$

- ▶  $S \subseteq \bigcup_{i=1}^n columns(R_i)$

- ▶  $p = p_1 \wedge \cdots \wedge p_m$

where each  $p_i$  is of the form “ $A \text{ op } c$ ” or “ $A \text{ op } A'$ ”

- ▶  $op$  is a comparison operator

- ▶  $c$  is a constant

- ▶  $A, A' \in \bigcup_{i=1}^n columns(R_i)$

# Eliminating Redundant DISTINCT

- ▶ Let  $Q = \pi_S^{all} \sigma_p(R_1 \times \cdots \times R_n)$  be a SPJ query
- ▶ Let  $F$  be the set of FDs defined as follows:
  - ▶ For each FD  $X \rightarrow Y$  that holds on  $R_1, \dots, R_n$ ,  $X \rightarrow Y \in F$
  - ▶ For each selection predicate  $A = c$  in  $p$ ,
$$\{A' \rightarrow A \mid A' \in \bigcup_{i=1}^n \text{columns}(R_i)\} \subseteq F$$
  - ▶ For each equality join predicate  $A = A'$  in  $p$ ,
$$\{A \rightarrow A', A' \rightarrow A\} \subseteq F$$

If  $S$  is a superkey of  $\sigma_p(R_1 \times \cdots \times R_n)$  wrt  $F$ , then the result of  $Q$  has no duplicates

# Example

- ▶ **R**(A,B,C)
- ▶ **S**(D,E,F)
- ▶ **T**(G,H)

```
select distinct B, D  
from   R, S, T  
where  A = E  
and    C = H  
and    G = 10
```



# Rewriting of Nested Queries

- ▶ Customer (cust#, cname, country)
- ▶ Order (order#, cust#, date, totalprice)

```
select  cname
from    Customer c
where    1000 <
          (select sum(o.totalprice)
           from    Order o
           where   o.cust# = c.cust#)
```

```
select    cname
from      Customer c join Order o
           on c.cust# = o.cust#
group by c.cust#, c.cname
having    1000 < sum(o.totalprice)
```

# Nested Queries

- ▶ A **nested query** is a query containing some subquery
- ▶ A subquery in a nested query is also called an **inner query** which is contained in an **outer query**
- ▶ A **correlated nested query** is a nested query where there is a subquery that is dependent on the tuple referenced in its outer query

# Subqueries in SQL

- ▶ Used as a scalar expression in SELECT clause

```
select order#,  
        (select c.cname  
         from   Customer c  
         where  c.cust# = o.cust#)  
from    Order o
```

# Subqueries in SQL (cont.)

- Used as a derived table in FROM clause

```
select c.cname
from    Customer c,
        ( select    cust#
          from      Order
          group by cust#
          having    1000 < sum(totalprice)
        ) as o
where   c.cust# = o.cust#
```

# Subqueries in SQL (cont.)

- ▶ Used as a scalar expression in WHERE / HAVING clause

```
select cname
from    Customer c
where    1000 <
          ( select sum(o.totalprice)
            from    Order o
            where    o.cust# = c.cust# )
```

# Classification of nested queries

- ▶ Is the nested query correlated or uncorrelated?
- ▶ Does the subquery involve aggregation?

# Classification of nested queries (cont.)

**Q1:** **select** order#  
**from** Order  
**where** cust# **in** (**select** cust#  
                  **from** Customer  
                  **where** country = "Singapore")

**Q2:** **select** order#  
**from** Order  
**where** totalprice = (**select** max(totalprice) **from** Order)

# Classification of nested queries (cont.)

**Q3:**

```
select cname
from   Customer c
where  exists (select *
               from   Order o
               where  o.cust# = c.cust#
               and    totalprice > 5000)
```

**Q4:**

```
select cname
from   Customer c
where  1000 <
      (select sum(o.totalprice)
       from   Order o
       where  o.cust# = c.cust#)
```



# How to unnest subqueries?

```
select  cname
from    Customer c
where   1000 <
        ( select sum(o.totalprice)
          from    Order o
          where   o.cust# = c.cust# )
```

```
select  cname
from    Customer c join Order o
        on c.cust# = o.cust#
group by c.cust#, c.cname
having   1000 < sum(o.totalprice)
```


- ▶ Transform subqueries into relational algebra expressions
  - ▶ Using a new relational operator: **apply** operator
- ▶ Apply rewriting rules to eliminate apply operator

# More Relational Algebra

$\mathcal{G}_{A,F} R$  = group by on relation R

- ▶ A = set of grouping columns
- ▶ F = set of aggregate functions

```
select   cname, sum(totalprice) as T
from     Customer c join Order o
           on c.cust# = o.cust#
group by cust#, cname
```




# More Relational Algebra (cont.)

$\mathcal{G}_F^1 R$  = scalar aggregation on relation R

- ▶  $F$  = set of aggregate functions
- ▶  $\mathcal{G}_F^1 R$  returns exactly one row

```
select order#  
from Order  
where totalprice >  
  (select sum(totalprice)  
   from Customer c join Order o  
   on c.cust# = o.cust#  
   where cname = "IBM")
```



# More Relational Algebra (cont.)

- ▶  $\mathcal{G}_{A,F} \emptyset = \emptyset$
- ▶  $\mathcal{G}_F^1 R$  returns exactly one row
- ▶ Thus,  $\mathcal{G}_{\emptyset,F} R \neq \mathcal{G}_F^1 R$

$$\mathcal{G}_F^1 R = \begin{cases} \mathcal{G}_{\emptyset,F} R & \text{if } R \neq \emptyset \\ (F_1(\emptyset), \dots, F_n(\emptyset)) & \text{otherwise} \end{cases}$$

where  $F = \{F_1, \dots, F_n\}$

# More Relational Algebra (cont.)

- If  $R$  is an empty relation, then

Query	Result
SELECT COUNT(A) FROM R	0
SELECT COUNT(*) FROM R	0
SELECT MIN(A) FROM R	NULL
SELECT MAX(A) FROM R	NULL
SELECT SUM(A) FROM R	NULL
SELECT AVG(A) FROM R	NULL

# Parameterized Relational Expression

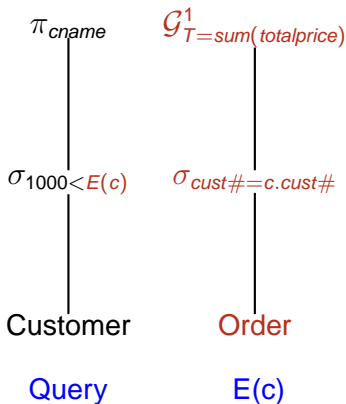
```
select cname  
from   Customer c  
where 1000 <  
      (select sum(o.totalprice)  
       from   Order o  
       where  o.cust# = c.cust#)
```

$E(c)$  = PRE with parameter  $c \in \textit{Customer}$

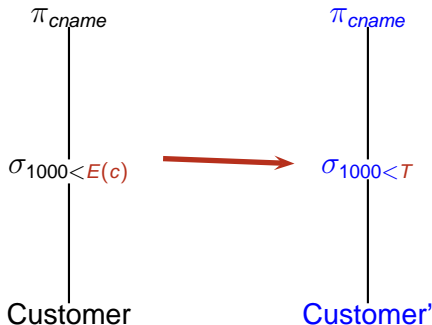
```
select cname  
from   Customer c  
where 1000 <  $E(c)$ 
```

# Parameterized Relational Expression

```
select cname
from   Customer c
where 1000 <
      (select sum(o.totalprice)
       from   Order o
       where  o.cust# = c.cust#)
```



# Eliminating PRE



- ▶ *Customer'* = *Customer* + additional column  $T$ 
  - ▶  $T = E(c)$
- ▶  $Customer' = \bigcup_{c \in Customer} (\{c\} \times E(c))$



# Eliminating PRE (cont.)

$$\text{Customer}' = \bigcup_{c \in \text{Customer}} (\{c\} \times E(c))$$

$$E(c) = \mathcal{G}_{T=\text{sum}(\text{totalprice})}^1(\sigma_{\text{cust}\# = c.\text{cust}\#}(\text{Order}))$$

**Customer**

cust#	cname	country
1	Alice	...
2	Bob	...
3	Carol	...

**Order**

order#	cust#	date	totalprice
100	2	...	3000
201	2	...	5000
460	3	...	4000
500	3	...	1000

**Customer'**

cust#	cname	country	T
1	Alice	...	null
2	Bob	...	8000
3	Carol	...	5000

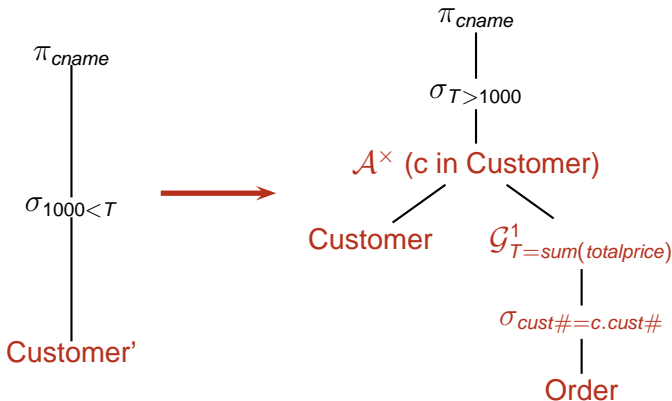
# Apply Operator $\mathcal{A}^\otimes$

$$R \mathcal{A}^\otimes E = \bigcup_{r \in R} (\{r\} \otimes E(r))$$

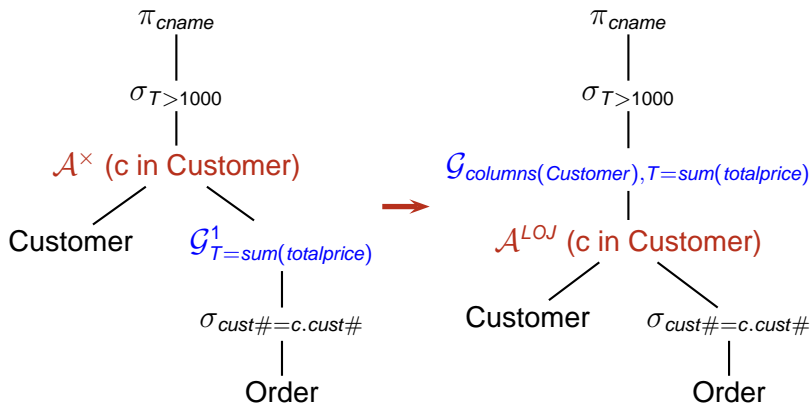
- ▶  $R$  = relational expression
- ▶  $E(r)$  = parameterized relational expression,  $r \in R$
- ▶  $\otimes$  = relational operator that combines  $R$  and  $E(r)$ 
  - ▶  $\otimes \in \{\times, \bowtie, \ltimes, \rtimes, \triangleright, \Join, \Joinr, \Joinl\}$
- ▶  $\bigcup$  = union all

# Eliminating PRE with apply operator

$$\begin{aligned}
 \text{Customer}' &= \bigcup_{c \in \text{Customer}} (\{c\} \times E(c)) \\
 &= \text{Customer } \mathcal{A}^\times E
 \end{aligned}$$



# Push down apply operator



**Rewriting rule:**  $R \mathcal{A}^\times (\mathcal{G}_F^1 E) = \mathcal{G}_{columns(R), F'} (R \mathcal{A}^{LOJ} E)$

# Example 1

Customer	
cust#	cname
1	Alice
2	Bob
3	Carol

Order		
order#	cust#	totalprice
100	2	3000
201	2	5000
460	3	4000
500	3	1000

Customer  $\mathcal{A}^{\times} \mathcal{G}_{T=\text{sum}(\text{totalprice})}^1 \sigma_{\text{cust\#}=c.\text{cust\#}}(\text{Order})$

cust#	cname	T
1	Alice	null
2	Bob	8000
3	Carol	5000

$R = \text{Customer } \mathcal{A}^{\text{LOJ}} \sigma_{\text{cust\#}=c.\text{cust\#}}(\text{Order})$

cust#	cname	order#	totalprice
1	Alice	null	null
2	Bob	100	3000
2	Bob	201	5000
3	Carol	460	4000
3	Carol	500	1000

$\pi_{\text{cust\#,cname,T}}(\mathcal{G}_{\text{columns}(\text{Customer}), T=\text{sum}(\text{totalprice})}(\mathbf{R}))$

cust#	cname	T
1	Alice	null
2	Bob	8000
3	Carol	5000



# Rewriting rule to push down apply

$$R \mathcal{A}^\times (\mathcal{G}_F^1 E) = \mathcal{G}_{columns(R), F'} (R \mathcal{A}^{LOJ} E)$$

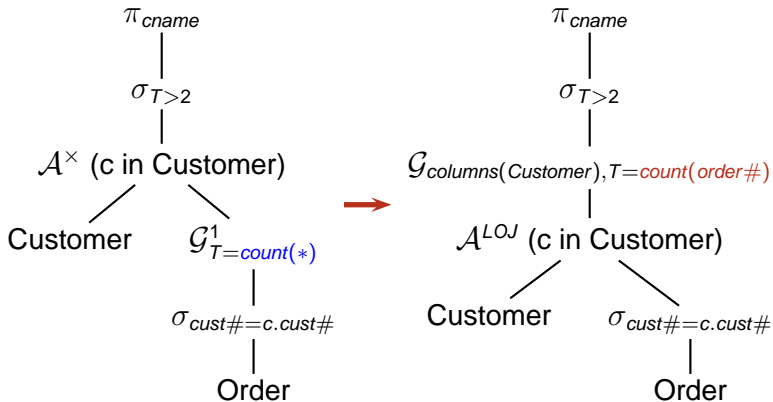
*iff*

- (1)  $R$  must contain a key, and
- (2)  $F'$  contains aggregates in  $F$  expressed over a single column

**Example:** If  $F$  is  $\text{COUNT}^*$ , then  $F'$  is  $\text{COUNT}(A)$  for some non-nullable column  $A$  from  $E$


## Example 2

**select** cname **from** Customer c **where**  $2 <$   
**(select count(\*) from** Order o **where** o.cust# = c.cust#)



# Example 2 (cont.)

Customer  $\mathcal{A}^\times \mathcal{G}_{T=\text{count}(*)}^1 \sigma_{\text{cust\#}=c.\text{cust\#}}(\text{Order})$



Customer	
cust#	cname
1	Alice
2	Bob

Order			
order#	cust#	date	totalprice
100	2	2012-05-03	3000
201	2	null	5000

cust#	cname	T
1	Alice	0
2	Bob	2

$R = \text{Customer } \mathcal{A}^{\text{LOJ}} \sigma_{\text{cust\#}=c.\text{cust\#}}(\text{Order})$

cust#	cname	order#	date	totalprice
1	Alice	null	null	null
2	Bob	100	2012-05-03	3000
2	Bob	201	null	5000

$\pi_{\text{cust\#}, \text{cname}, T}(\mathcal{G}_{\text{columns}(\text{Customer}), T=\text{count}(\text{order\#})}(\mathbf{R}))$

cust#	cname	T
1	Alice	0
2	Bob	2

$\pi_{\text{cust\#}, \text{cname}, T}(\mathcal{G}_{\text{columns}(\text{Customer}), T=\text{count}(*)}(\mathbf{R}))$

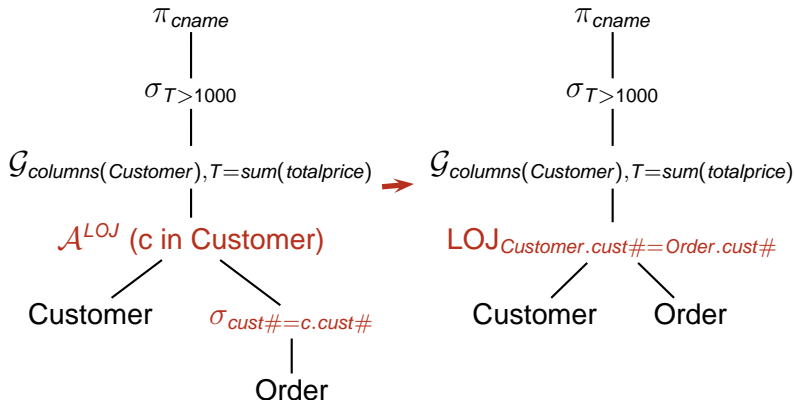
cust#	cname	T
1	Alice	1
2	Bob	2

$\pi_{\text{cust\#}, \text{cname}, T}(\mathcal{G}_{\text{columns}(\text{Customer}), T=\text{count}(\text{date})}(\mathbf{R}))$

cust#	cname	T
1	Alice	0
2	Bob	1

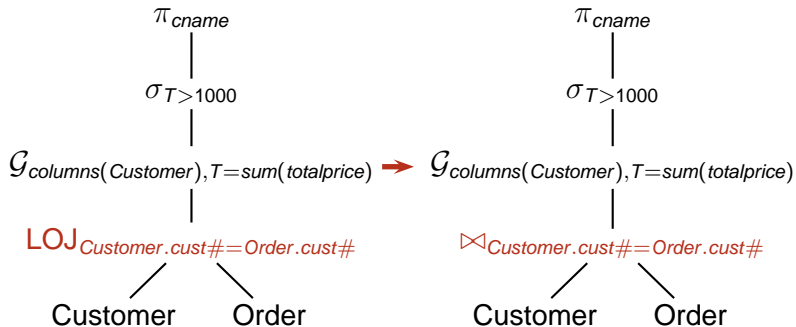


# Eliminate apply operator



**Rewriting rule:**  $R \mathcal{A}^{\otimes} (\sigma_p E) = R \otimes_p E$   
 $E$  must not have any parameter that refers to  $R$

# Simplify LOJ to inner join



# Therefore ...

```
select  cname
from    Customer c
where    1000 <
          (select sum(o.totalprice)
           from    Order o
           where   o.cust# = c.cust#)
```

```
select  cname
from    (select    cname, sum(totalprice) as T
          from      Customer c join Order o
                  on c.cust# = o.cust#
          group by cust#, cname)
where    1000 < T
```

# Therefore ...

```
select  cname
from    Customer c
where    1000 <
          ( select sum(o.totalprice)
            from    Order o
            where   o.cust# = c.cust#)
```

```
select    cname
from      Customer c join Order o
          on c.cust# = o.cust#
group by c.cust#, c.cname
having    1000 < sum(o.totalprice)
```

# Rewriting rules for apply operator

1.  $R \mathcal{A}^{\otimes} E = R \otimes_{true} E$
2.  $R \mathcal{A}^{\otimes} (\sigma_p E) = R \otimes_p E$
3.  $R \mathcal{A}^{\times} (\sigma_p E) = \sigma_p (R \mathcal{A}^{\times} E)$
4.  $R \mathcal{A}^{\times} (\pi_v E) = \pi_{v \cup columns(R)} (R \mathcal{A}^{\times} E)$
5.  $R \mathcal{A}^{\times} (E_1 \cup E_2) = (R \mathcal{A}^{\times} E_1) \cup (R \mathcal{A}^{\times} E_2)$
6.  $R \mathcal{A}^{\times} (E_1 - E_2) = (R \mathcal{A}^{\times} E_1) - (R \mathcal{A}^{\times} E_2)$

- Rules 1 & 2 require that  $E$  must not have any parameter that refers to  $R$

# Rewriting rules for apply operator (cont.)

$$7. R \mathcal{A}^\times (E_1 \times E_2) = (R \mathcal{A}^\times E_1) \bowtie_{R.key} (R \mathcal{A}^\times E_2)$$

$$8. R \mathcal{A}^\times (\mathcal{G}_{A,F} E) = \mathcal{G}_{A \cup \text{columns}(R), F} (R \mathcal{A}^\times E)$$

$$9. R \mathcal{A}^\times (\mathcal{G}_F^1 E) = \mathcal{G}_{\text{columns}(R), F'} (R \mathcal{A}^{LOJ} E)$$

- ▶ Rules 7 to 9 require that  $R$  contains a key (denoted by  $R.key$ )
- ▶ In Rule 7,  $\bowtie_{R.key}$  denote an equality join on  $R.key$

# References

## Required Readings

- ▶ Section 4.6 (Query Tuning) of Shasha & Bonnet's book
- ▶ C.A. Galindo-Legaria, M.M. Joshi, *Orthogonal optimization of subqueries and aggregation*, SIGMOD 2001