

# **CS5226 Lecture 4**

## **Memory Tuning**

# Memory Management


- ▶ **Shared data and control structures:**

- ▶ buffer cache
- ▶ catalog metadata
- ▶ log buffer
- ▶ lock structures
- ▶ SQL execution plans
- ▶ etc

- ▶ **Local data and control structures:**

- ▶ **operator work area** = memory allocated for evaluating an SQL operator
- ▶ Memory-intensive operators:
  - ★ sort
  - ★ hash join

# Oracle Memory Model



(Dageville & Zait, VLDB 2002)

**System Global Area (SGA)** = shared memory region

**Process Global Area (PGA)** = private memory region

# Oracle Memory Management

- ▶ **Pre 9i - Manual memory mangement**
  - ▶ **Parameters for tuning SGA:**
    - ★ DB\_CACHE\_SIZE, SHARED\_POOL\_SIZE, LARGE\_POOL\_SIZE, JAVA\_POOL\_SIZE, etc.
  - ▶ **Parameters for tuning PGA:**
    - ★ SORT\_AREA\_SIZE, HASH\_AREA\_SIZE, BITMAP\_MERGE\_AREA\_SIZE, etc.
- ▶ **9i - Automatic PGA memory management**
  - ▶ PGA\_AGGREGATE\_TARGET parameter
- ▶ **10g - Automatic SGA memory management**
  - ▶ SGA\_TARGET parameter
- ▶ **11g - Automatic memory management**
  - ▶ MEMORY\_TARGET parameter

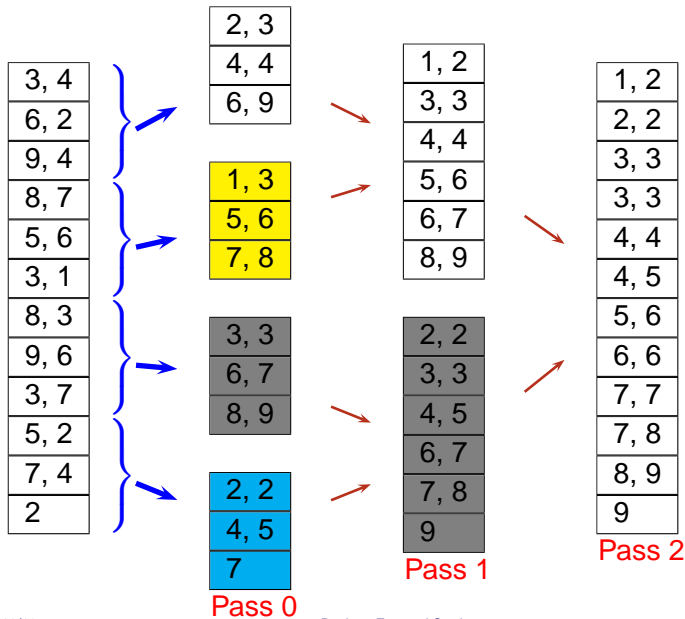
# Memory-Intensive Operators

- ▶ Examples of memory-intensive operators:
  - ▶ sort
  - ▶ hash join
- ▶ How to manage memory for such operators?

# Review: External Merge Sort

- ▶ Sorting file of  $N$  pages with  $B$  buffer pages
- ▶ **Pass 0: Creation of sorted runs**
  - ▶ Read in and sort  $B$  pages at a time
  - ▶ Number of sorted runs created =  $\lceil N/B \rceil$
  - ▶ Size of each sorted run =  $B$  pages (except possibly for last run)
- ▶ **Pass  $i$ ,  $i \geq 1$ : Merging of sorted runs**
  - ▶ Use  $B - 1$  buffer pages for input & one buffer page for output
  - ▶ Performs  $(B-1)$ -way merge
- ▶ **Analysis:**
  - ▶  $N_0$  = number of initial sorted runs =  $\lceil N/B \rceil$
  - ▶ total number of passes =  $\lceil \log_{B-1}(N_0) \rceil + 1$

# External Merge Sort with B=3: Example



# Review: Hash Join, $R \bowtie S$

- ▶ Idea:

- ▶ Partition  $R$  and  $S$  into  $k$  partitions using some hash function  $h$

- ★  $R = R_1 \cup R_2 \cup \dots \cup R_k$

- ★  $S = S_1 \cup S_2 \cup \dots \cup S_k$

- ▶ Joins corresponding pair of partitions

- ★  $R \bowtie S = (R_1 \bowtie S_1) \cup (R_2 \bowtie S_2) \cup \dots \cup (R_k \bowtie S_k)$

- ▶ Algorithms:

- ▶ Grace hash join
  - ▶ Hybrid hash join



# Review: Grace Hash Join $R \bowtie S$

- ▶ Consists of three phases:
  1. Partition  $R$  into  $R_1, \dots, R_k$
  2. Partition  $S$  into  $S_1, \dots, S_k$
  3. Probing phase: probes each  $R_i$  with  $S_j$ 
    - ★ Read  $R_i$  to build a hash table
    - ★ Read  $S_j$  to probe hash table
- ▶  $R$  is called the **build relation** &  $S$  is called the **probe relation**

## Partitioning (building) phases

initialize a hash table  $T$  with  $k$  buckets

for each tuple  $r \in R$  do

    insert  $r$  into bucket  $h(r)$  of  $T$

write each bucket  $R_i$  of  $T$  to disk

initialize a hash table  $T$  with  $k$  buckets

for each tuple  $s \in S$  do

    insert  $s$  into bucket  $h(s)$  of  $T$

write each bucket  $S_i$  of  $T$  to disk

## Probing (matching) phase

for  $i = 1$  to  $k$  do

    initialize a hash table  $T$

    for each tuple  $r$  in partition  $R_i$  do

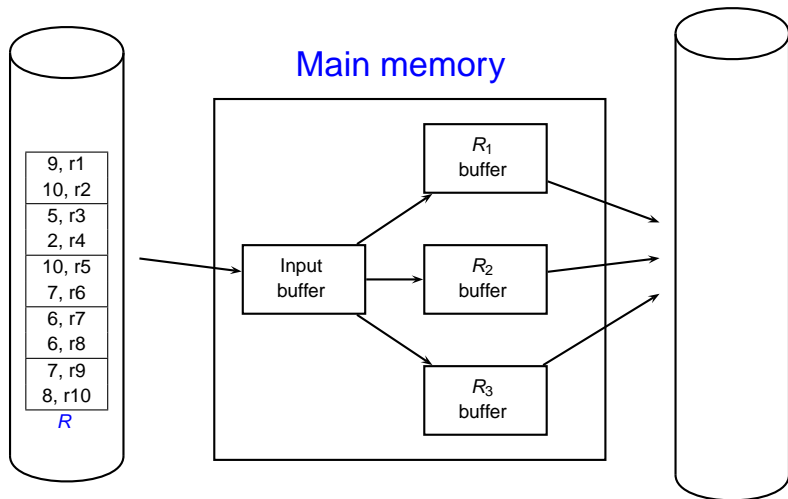
        insert  $r$  into bucket  $h'(r)$  of  $T$

    for each tuple  $s$  in partition  $S_i$  do

        for each tuple  $r$  in bucket  $h'(s)$  of  $T$  do


            if  $r$  and  $s$  matches then output  $(r, s)$

# Grace Hash Join: Partitioning Relation R




Hash function:  $h(v) = v \bmod 3$


# Grace Hash Join: Partitioning Relation R




# Grace Hash Join: Partitioning Relation R




# Grace Hash Join: Partitioning Relation R




# Grace Hash Join: Partitioning Relation R



# Grace Hash Join: Partitioning Relation R




# Grace Hash Join: Partitioning Relation R







# Grace Hash Join: Partitioning Relation R



# Grace Hash Join: Partitioning Relation R




# Grace Hash Join: Partitioning Relation R




$R$  is partitioned into  $R_1$ ,  $R_2$ , &  $R_3$

# Grace Hash Join: Partitioning Relation S




Similarly,  $S$  is partitioned into  $S_1$ ,  $S_2$ , &  $S_3$


# Grace Hash Join: Probing Phase




# Grace Hash Join: Probing Phase



# Grace Hash Join: Probing Phase




# Grace Hash Join: Probing Phase







# Grace Hash Join: Probing Phase




# Grace Hash Join: Probing Phase




# Grace Hash Join: Probing Phase




# Grace Hash Join: Probing Phase




# Grace Hash Join: Probing Phase



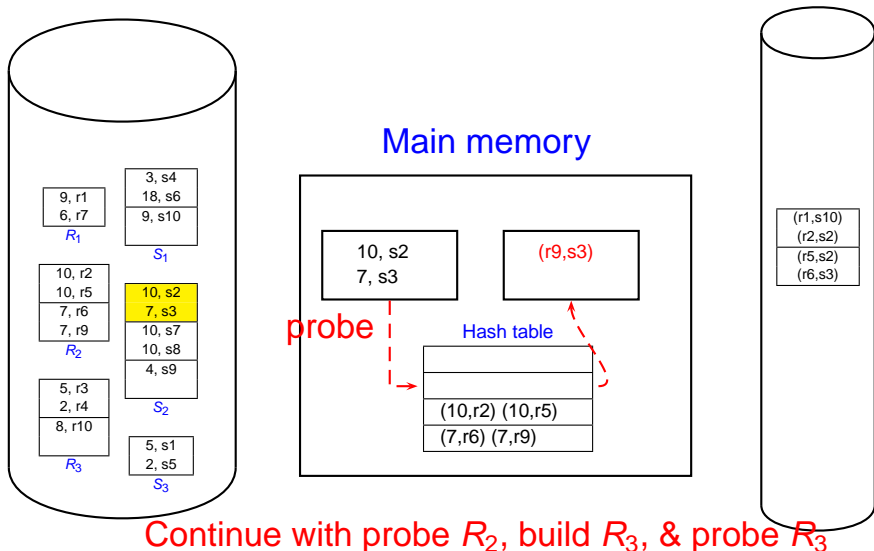
# Grace Hash Join: Probing Phase



# Grace Hash Join: Probing Phase



# Grace Hash Join: Probing Phase





# Review: Hybrid Hash Join (HHJ), $R \bowtie S$

- ▶ Improvement of Grace Hash Join
- ▶ Enables some build partitions  $R_i$  to be resident in memory at the end of **partitioning build relation  $R$** 
  - ▶ Each  $R_i$  is either a **resident** or **spilled** partition
  - ▶ Memory usage
    - ★ hash tables for resident partitions
    - ★ output buffers for spilled partitions
    - ★ input buffer
    - ★ free buffer - spool area for pages flushed to disk
- ▶ During **partitioning of probe table  $S$** ,  $S_i$  can join with  $R_i$  if  $R_i$  is resident partition
- ▶ Number of partitions,  $k = \sqrt{F|R|}$ 
  - ▶  $|R|$  = size of build relation  $R$  in pages
  - ▶  $F$  = fudge factor

# HHJ: Partition Build Relation

01. let  $k = \sqrt{F|R|}$
02. for  $i = 1$  to  $k$  do
03.     mark  $R_i$  as a resident partition & allocate one page to  $R_i$
04.     allocate input buffer & assign remaining pages to free buffer
05.     for each tuple  $t \in R$  do
06.         let  $t$  be hashed to  $R_i$
07.         if ( $R_i$  is full) and ( $R_i$  is resident) then
08.             if (there's a free page) then
09.                 add free page to  $R_i$
10.             else
11.                 mark  $R_i$  as spilled
12.                 flush  $R_i$  to disk & free all pages (except one) of  $R_i$
13.             else if ( $R_i$  is full) and ( $R_i$  is spilled) then
14.                 flush  $R_i$  to disk
15.             insert  $t$  into  $R_i$
16.     flush & free all pages in spilled partitions

# HHJ: Partition Probe Relation

01. for  $i = 1$  to  $k$  do
02.     if ( $R_i$  is spilled) then
03.         allocate one page to  $S_i$
04.     for each tuple  $t \in S$  do
05.         let  $t$  be hashed to  $S_i$
06.         if ( $R_i$  is resident) then
07.             use  $t$  to probe  $R_i$  for matches
08.         else
09.             insert  $t$  into  $S_i$
10.             if ( $S_i$  is full) then
11.                 flush  $S_i$  to disk
12. flush all pages in  $S$  partitions & free all pages

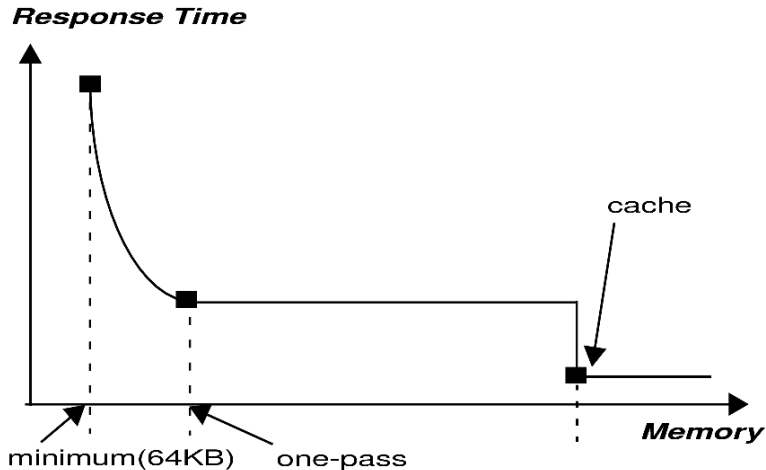
# HHJ: Probing Phase

01. for  $i = 1$  to  $k$  do
02.     if ( $R_i$  is spilled) then
03.         read in  $R_i$  & build a hash table
04.         for each tuple  $t \in S_i$  do
05.             use  $t$  to probe  $R_i$  for matches

# Oracle PGA Memory Management


- ▶ **operator work area** = memory allocated for operator's execution
- ▶ Oracle's classification of **work area size**:
  - ▶ **cache size** = minimum work area size that enables evaluation to incur a single scan of input data
    - ★  $\text{cache size} = \text{size of input data} + \text{auxiliary memory structures}$
  - ▶ **one-pass size** = minimum work area size that enables evaluation to incur 2 scans of input data
  - ▶ **multiple-pass size** = work area size that is less than one-pass size
- ▶  $\text{cache size} > \text{one-pass size} > \text{multiple-pass size}$

# Effect of Memory on Sort Performance



(Dageville & Zait, VLDB 2002)

# Effect of Memory on Hash Join Perf.



(Dageville & Zait, VLDB 2002)

# Memory Management Approaches

How much memory to allocate to SQL operators?

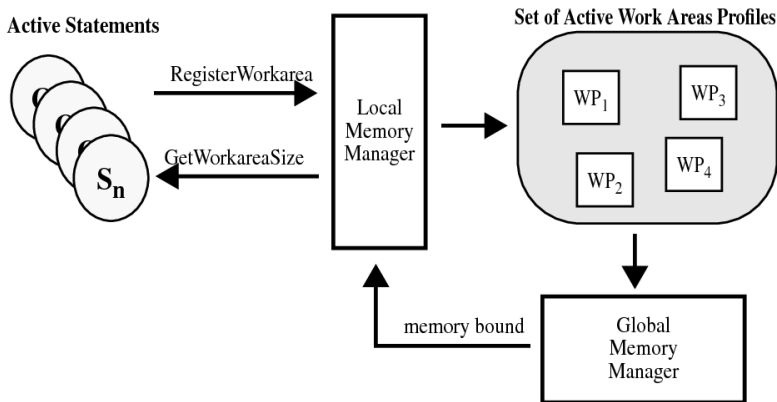
- ▶ Fixed amount of memory for each operator
- ▶ Amount based on size of input operands
- ▶ Amount based on size of input operands + current workload
- ▶ Memory usage adapts to memory demand of system



# DBMS Approaches (circa 2002)

DBMS	Initial work area size	Size during execution	Response to memory pressure
Oracle 8i	static	static	none
<b>Oracle 9i</b>	<b>dynamic</b>	<b>adaptive</b>	<b>adaptive</b>
DB2/UDB 7.1	static	static	minimum
Informix 9.3	static	static	limit ops
SQL Server 7	dynamic	static	queueing
Teradata	dynamic	static	?

# Memory Management Feedback Loop



(Dageville & Zait, VLDB 2002)

# Work Area Profile

**Work area profile** = metadata for an **operator work area**:

- ▶ Operator type
  - ▶ Example: sort, hash-join
- ▶ Current memory requirement to run with minimum/one-pass/cache memory
- ▶ Number instances of work area
  - ▶ ie degree of parallelism of operator
- ▶ Current amount of PGA memory

# Memory Management Feedback Loop

## SQL Operator

1. Registers **work area profile** with local memory manager
5. Adjusts memory usage based on **expected work area size** (derived by LMM) & updates **work area profile**

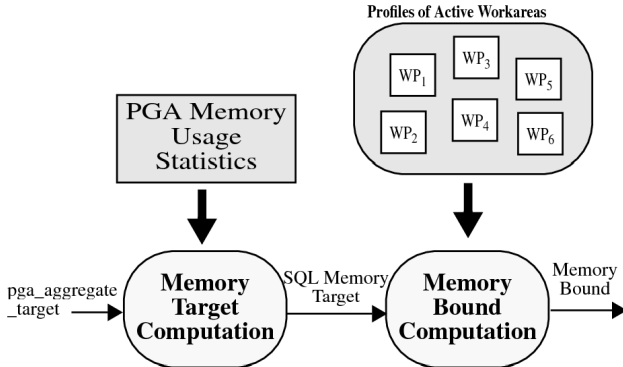
## Local Memory Manager (LMM)

2. Maintains **active work area profiles** in SGA
4. Computes **expected work area size** for each operator based on
  - ▶ **operator's work area profile**
  - ▶ **memory bound** (derived by GMM)

## Global Memory Manager (GMM)

3. Computes **memory bound** for work area based on
  - ▶ `pga_aggregated_target`
  - ▶ PGA memory usage statistics
  - ▶ **active work area profiles**

# GMM: Memory Bound Computation



(Dageville & Zait, VLDB 2002)

- ▶ **GMM** computes **memory bound** for work area based on
  - ▶ `pga_aggregated_target`
  - ▶ PGA memory usage statistics
  - ▶ **active work area profiles**
- ▶ **SQL memory target** = amount of PGA allocated to work areas

# Memory Bound Computation

- ▶ Find largest value of  $B$  such that

$$\sum_{i=1}^N \text{ExpectedWorkAreaSize}(WP_i, B) \leq \text{SQL Memory Target}$$

- ▶  $N$  = number of work area profiles
- ▶  $WP_i$  = memory profile of  $i^{\text{th}}$  work area,  $i \in [1, N]$
- ▶ Work area profile
  - ▶ minimum, one-pass & cache memory thresholds of work area

# Expected Work Area Size, $E$

- ▶ Computation of  $E$  depends on type of operator (sort or non-sort) & execution mode (serial or parallel)

- ▶  $E \in [\text{minimum threshold}, \text{cache threshold}]$

- ▶ **Type of operator:**

- ▶ Operator is sort

$$E = \begin{cases} \text{cache threshold} & \text{if } B \geq \text{cache threshold,} \\ \text{one-pass threshold} & \text{if } B \geq \text{one-pass threshold,} \\ \text{multi-pass threshold} & \text{otherwise.} \end{cases}$$

- ▶ Operator is non-sort

$$E = \min\{B, \text{cache threshold}\}$$


- ▶ **Execution mode:**

$$E = \begin{cases} \min\{E, 0.05 \times \text{memory target}\} & \text{if serial,} \\ \min\{E \times DOP, 0.3 \times \text{memory target}\} & \text{otherwise.} \end{cases}$$

# Example of Memory Bound Computation

(Dageville & Zait, VLDB 2002)

**Target = 133MB => Bound = 20M**






# Drift Management

- ▶ **GMM** is a background daemon
  - ▶ Recomputes **memory bound** periodically
- ▶ Limitation of background recomputation:
  - ▶ Slow to react to workload changes
- ▶ **Memory drift**
  - ▶ Measures staleness of memory bound
  - ▶ Expected amount of memory acquired/released by queries since last recomputation of memory bound
  - ▶ Drift can be positive or negative
- ▶ Triggers **foreground recomputation** when drift exceeds 10% of SQL memory target

# Foreground Computation of Bound



(Dageville & Zait, VLDB 2002)

# Memory Adaptive External Sorting

- ▶ **Creation of initial sorted runs**
  - ▶ Adapt size of sorted runs
- ▶ **Merging sorted runs**
  - ▶ Adapt merging fan-out
- ▶ Adapt size of input/output buffers

# Memory Adaptive Hash Join

- ▶ **Partitioning of build relation**
  - ▶ Adapt size of free buffer
  - ▶ Convert resident partitions to spilled ones to release memory
- ▶ **Partitioning of probe relation**
  - ▶ Adapt size of free buffer
  - ▶ Convert resident partitions to spilled ones to release memory
  - ▶ Restore spilled partitions to resident ones using extra memory
- ▶ Adapt size of input/output buffers

# References

## Required Readings

- ▶ B. Dageville, M. Zait, SQL Memory Management in Oracle9i, VLDB 2002, 962-973.

## Additional Readings

- ▶ W. Zhang, P.A. Larson, *Dynamic Memory Adjustment for External Mergesort*, VLDB 1997.
- ▶ H.H. Pang, M.J. Carey, M. Livny, *Partially preemptible hash joins*, SIGMOD 1993.
- ▶ A.J. Storm, C. Garcia-Arellano, S.S. Lightstone, Y. Diao, M. Surendra, *Adaptive self-tuning memory in DB2*, VLDB 2006.